


Generative Adversarial Networks



Generative models



- 1) Parametric model
 - 2) Non-parametric model
↳ histogram, kernel
 - 3) VAE
 - 4) GAN $\rightarrow p(z)$
- $p(x)$
↓


Generate new samples from the distribution of training data

Do not learn just the same samples as that of training data





desired data distribution

~



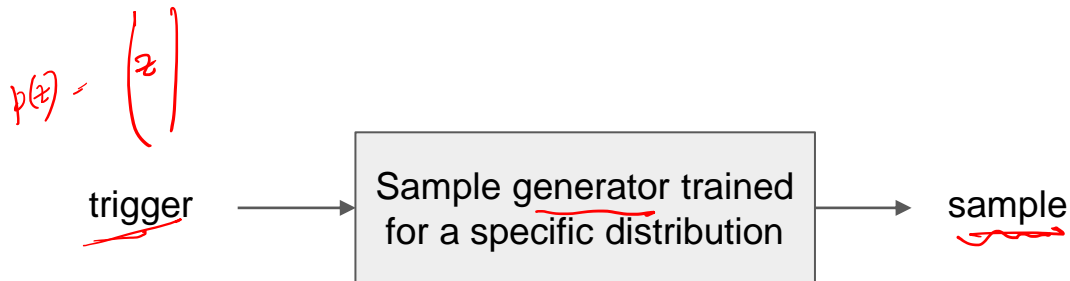
generated distribution



Generative Adversarial Network (GAN)

Real distributions are difficult to model as explicit density (by formula)

Instead, generate samples from such complex high-dimensional distributions



Generative Adversarial Network (GAN) does not model explicit density

Generative models

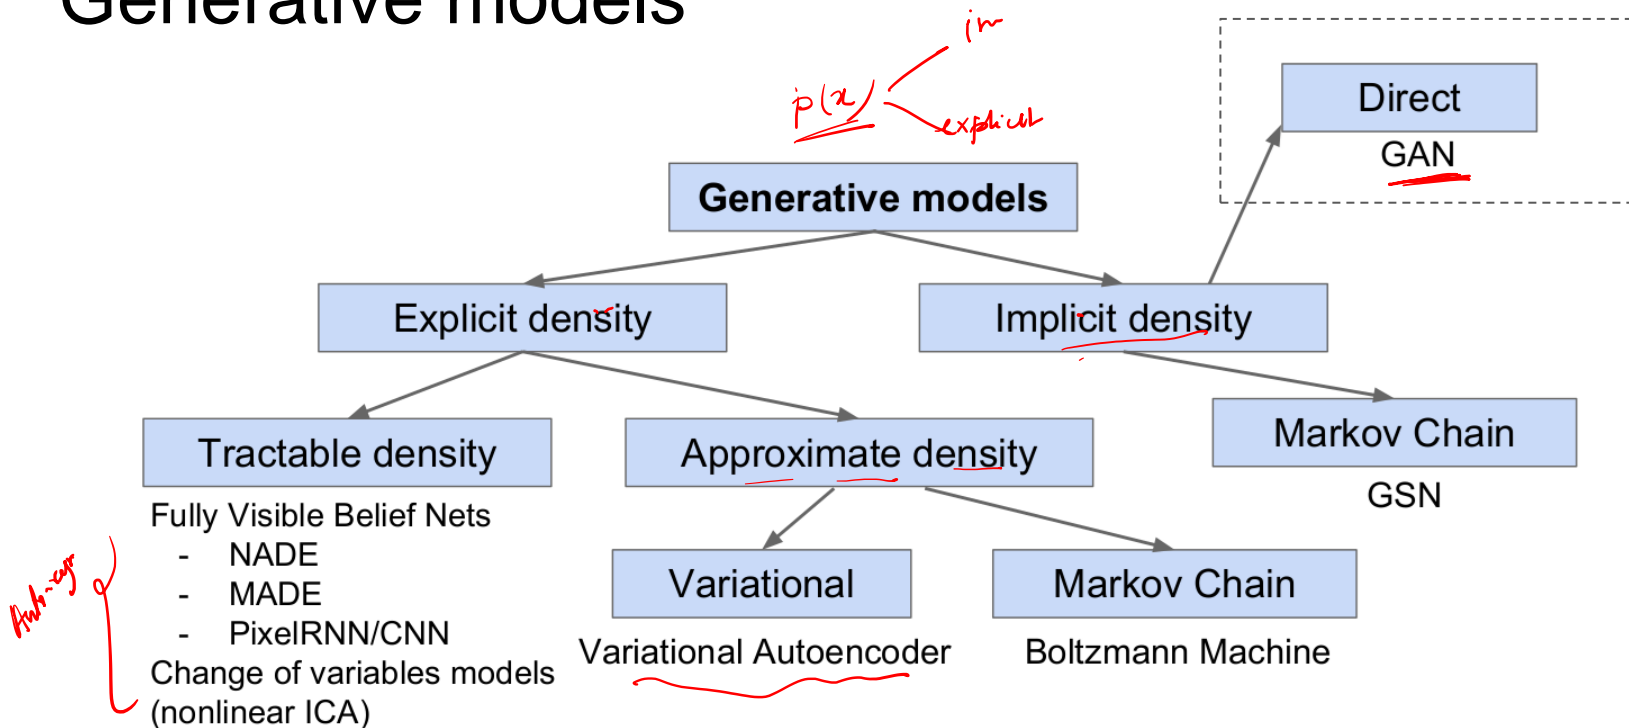



Figure copyright: CS231n 2017 Stanford course which is adapted from Ian Goodfellow tutorial on GANs 2017.

GAN

Generative Adversarial Networks

The Generator

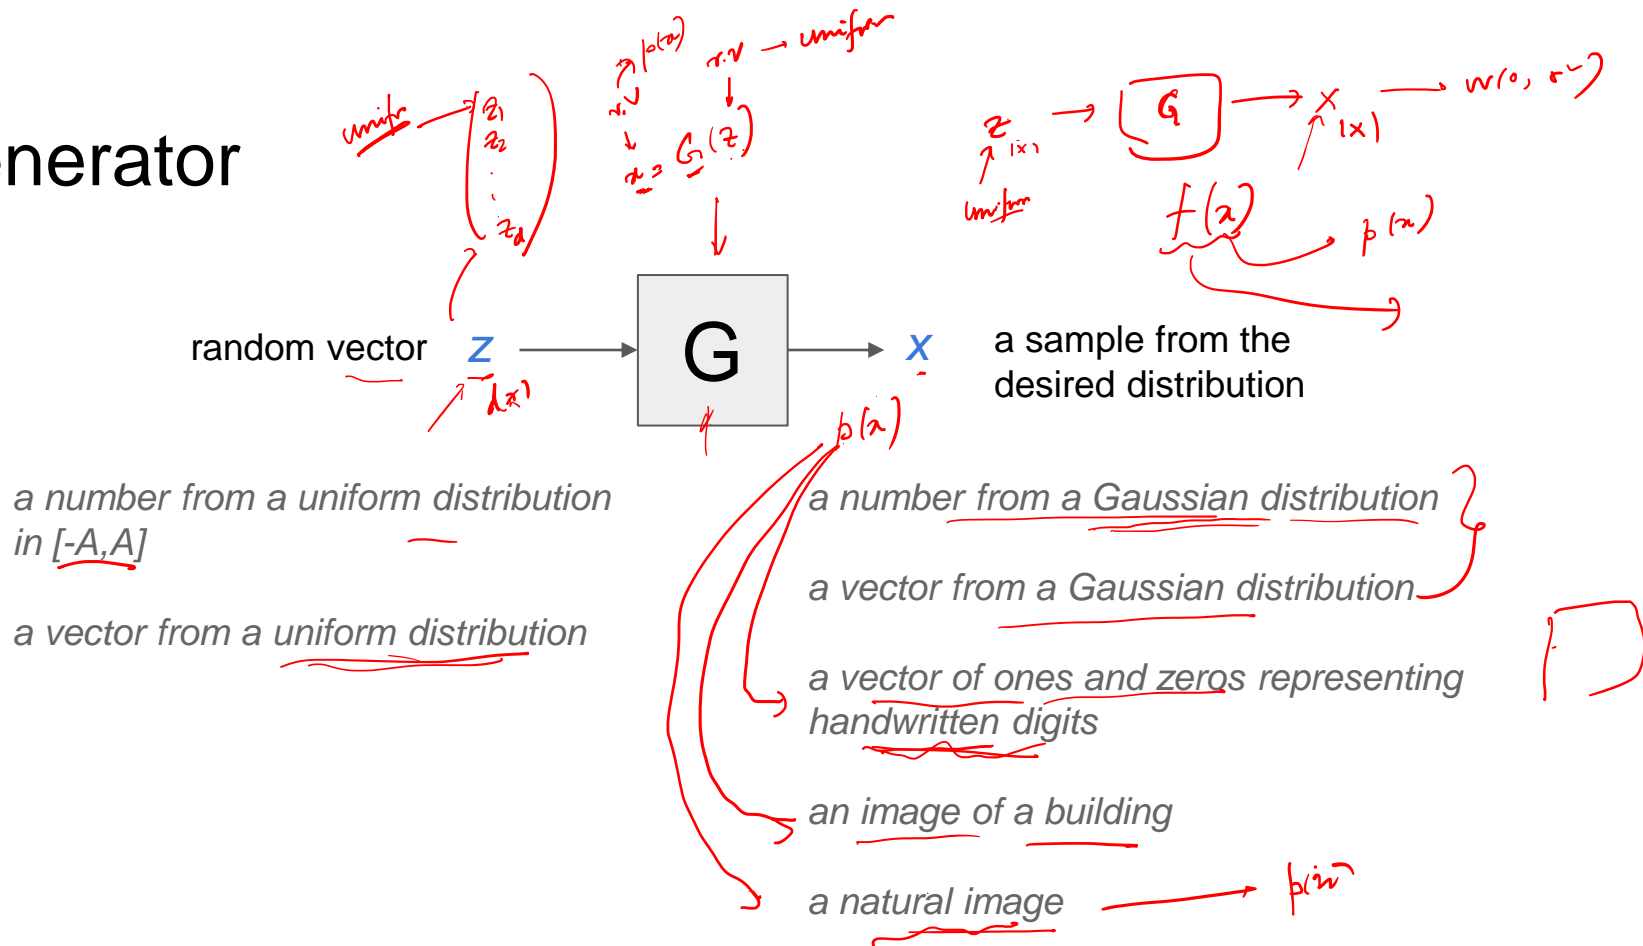


The goal is to create a system known as Generator G which will generate “new” samples from a distribution of data

The generator is trained to learn the distribution of data by its own, and is not given the actual distribution (by formula)


To generate different samples, it takes a random vector as its input (trigger)

Generator



The generator creates a mapping between these two spaces.

Generator




Adversarial idea

The generator should not generate the exact same samples used as training data

adversarial adj.- involving or characterized by conflict or opposition

Instead of training only the generator, we will train an additional system that will tell whether the generator generates samples belonging to the desired distribution, and not whether the generator exactly generates the same training data

Discriminator




D outputs 1 if the sample is from the data distribution \mathcal{D}

- it says the sample is "real"

D outputs 0 if the sample is from the generated distribution \mathcal{G}

- it says the sample is "fake"

Two player game




G tries to make $D(G(z))$ near 1
-- "I am not fake"



D tries to make $D(G(z))$ near 0
-- "You are fake"

Handwritten red note: $D(G(z)) \neq 1$



D tries to output 1
-- "You are real"

Adversarial framework

$0 \leq D(x) \leq 1$
 $\max_D \mathbb{E}_{x \sim p_{\text{data}}} [\log(D(x))]$
 Discriminator should classify real as 1
 $D(x) \approx 1$
 $\{x_1, \dots, x_n\}$
 $p_{\text{data}}(x)$
 $\max_D \mathbb{E}_{x \sim p_{\text{data}}} [\log(D(x))]$
 D should maximize $\log D(x)$ if x comes from data ("real")
 $\max_D \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$
 D should maximize $\log(1 - D(x))$ if x comes from generator ("fake")
 $\min_G \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$
 G should minimize $\log(1 - D(x))$ if x comes from generator (G competes by saying "I am real")
 $\mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \approx 0$

$$\min_G \max_D \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

$V(D, G)$

Training

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$. ↖ $G(z^{(i)})$
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$. ← real data
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Training

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient: D says generated $G(z)$ is fake, i.e. push $D(G(z))$ to 0

D says data x is real, i.e. push $D(x)$ to 1 $\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right]$.

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)})))$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

UPDATE
DISCRIMINATOR

Training

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by **descending** its stochastic gradient:

UPDATE
GENERATOR

minimize

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)})))$$

G says generated $G(z)$ is real, i.e. push $D(G(z))$ to 1

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Training

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by **descending** its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)})))$$

G says generated $G(z)$ is real, i.e. push $\log(1 - D(G(z)))$ to -inf

minimize



end for

The gradient-based updates can use any standard gradient descent in our experiments.

Decrease the log probability that the discriminator makes the correct prediction that "gen=fake"

Early in learning, when G is poor, D can reject samples with high confidence because they are clearly different from the training data. In this case, $\log(1 - D(G(z)))$ saturates.

$$\log(1 - D(G(z))) \approx 0$$

Training


Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$



The generator aims to increase the log probability that the discriminator makes a mistake,
i.e. do $\max \log D(G(z))$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by **ascending** its stochastic gradient:

maximize

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left(D(G(z^{(i)})) \right).$$

rather than aiming to decrease the log probability that the discriminator makes the correct prediction,
i.e. don't use $\min \log (1 - D(G(z)))$

G says generated $G(z)$ is real, i.e. push $\log D(G(z))$ to 0

end for

The gradient-based updates can use any standard gradient descent algorithm in our experiments.

Increase the log probability that the discriminator makes the mistake that "gen=real"

Training

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

UPDATE
DISCRIMINATOR
$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by ascending its stochastic gradient:

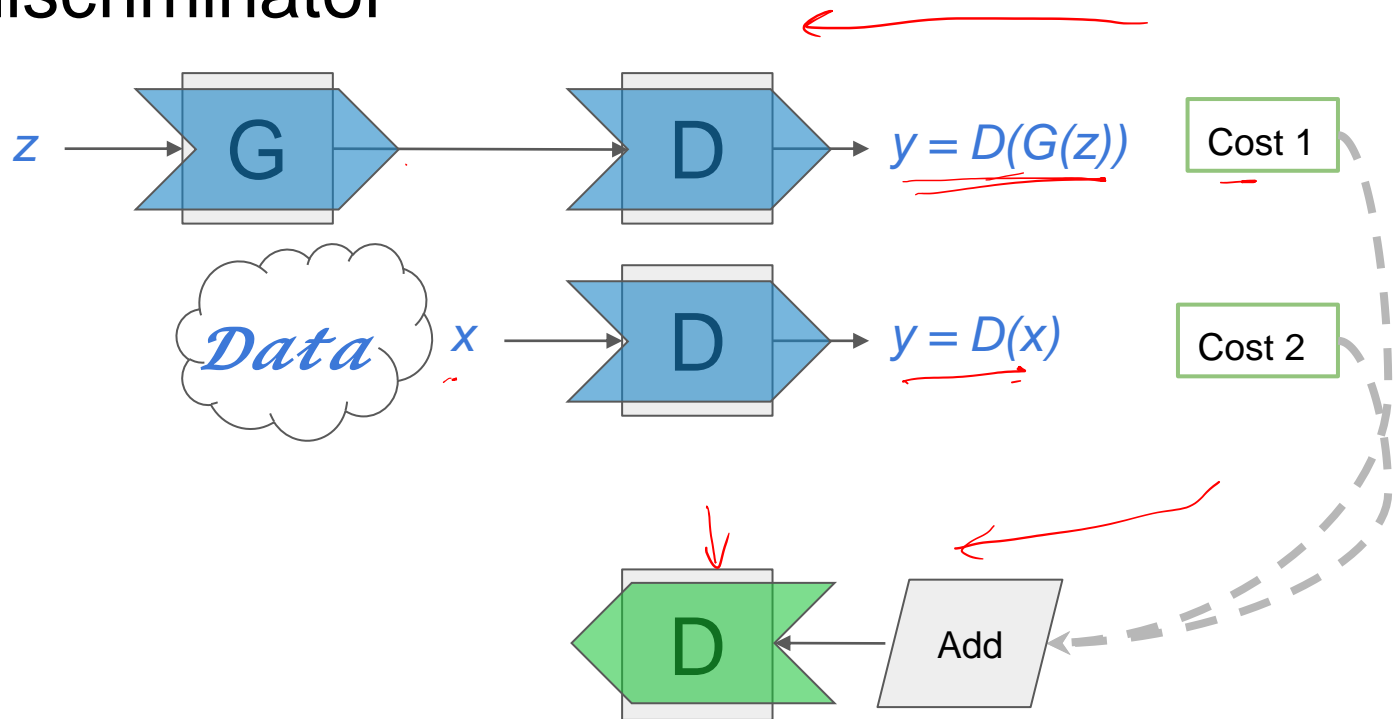
UPDATE
GENERATOR
$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (D(G(z^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

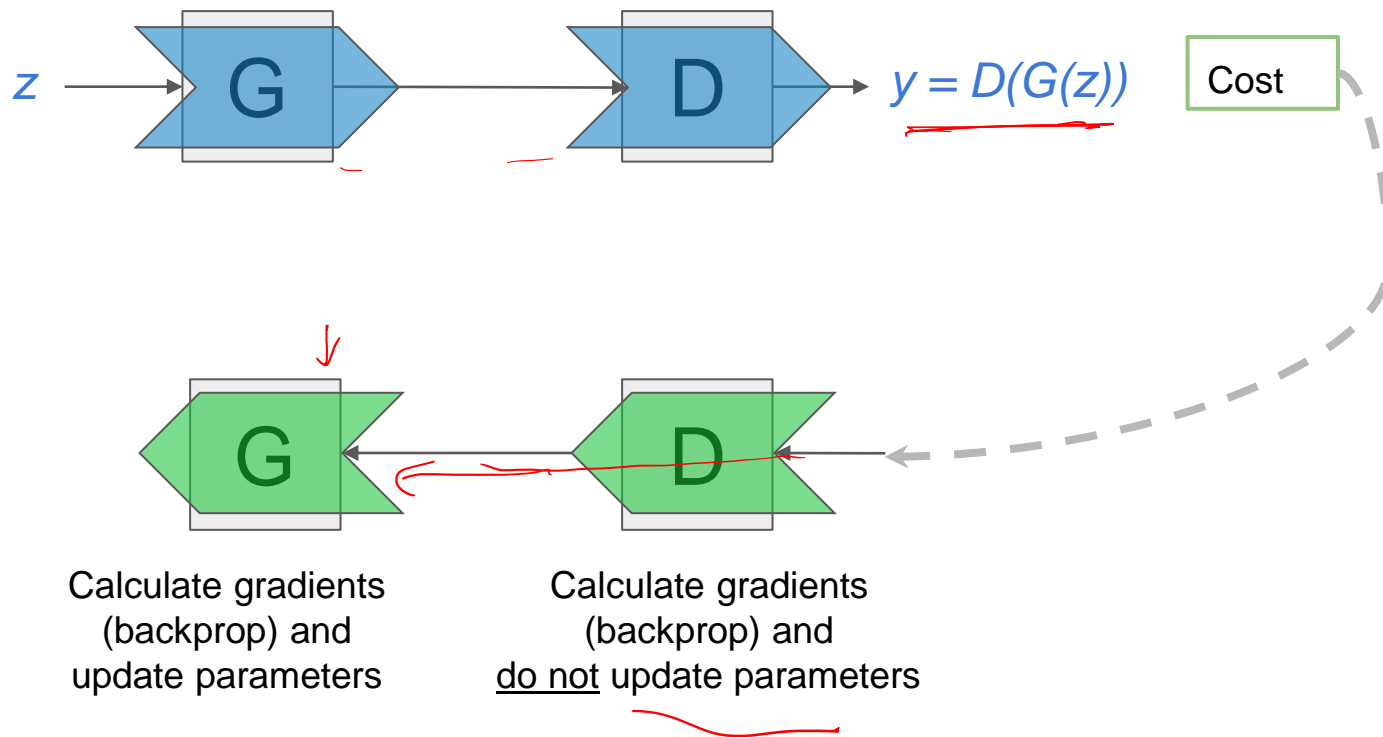
LOOP
UNTIL YOU ARE
SATISFIED WITH THE
GENERATOR SAMPLES

Update discriminator



Calculate gradients (backprop)
and update parameters

Update generator



Training example illustration

z is sampled from a uniform distribution, and data is Gaussian

Both z and x are single scalars

G and D are multilayer perceptrons


training both G and D

Legend

Data distribution (desired)

Generated distribution

Discriminator output



G maps z to x

Training example illustration

z is sampled from a uniform distribution, and data is Gaussian

Both z and x are single scalars

G and D are multilayer perceptrons


training both G and D

Legend

Data distribution (desired)

Generated distribution

Discriminator output



G maps z to x

G : green (generated dist) slowly moves towards black (data dist)

Training example illustration

z is sampled from a uniform distribution, and data is Gaussian

Both z and x are single scalars

G and D are multilayer perceptrons


Legend

Data distribution (desired)

Generated distribution

Discriminator output

training both G and D



G maps z to x

G : green (generated dist) slowly moves towards black (data dist)

D : initially, D (blue) classifies green as fake (low), and black as real (high).

Training example illustration

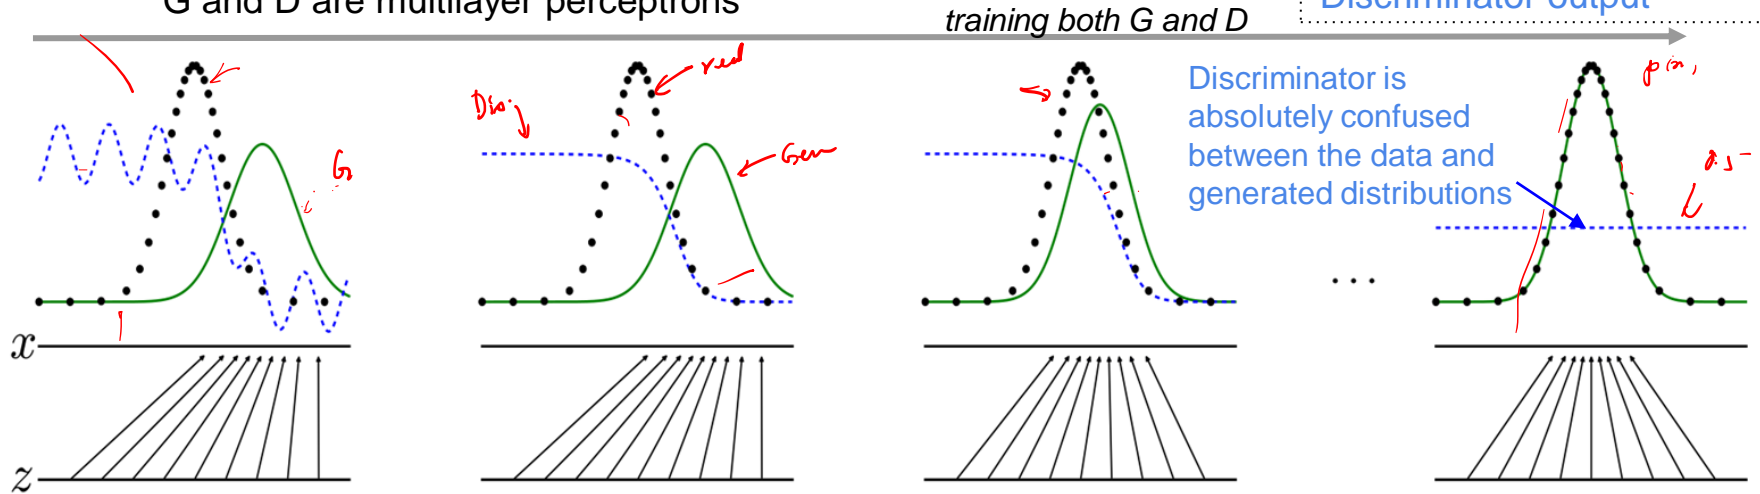
z is sampled from a uniform distribution, and data is Gaussian
Both z and x are single scalars
 G and D are multilayer perceptrons

Legend

Data distribution (desired)

Generated distribution

Discriminator output




At last, G produces samples from the distribution same as or close to the data distribution

- D finally outputs the same (0.5) for both generated and data distributions.
- **Important:** D does not try to identify the samples from the data distribution. D tries only to differentiate the data distribution and the generated distribution.

1D GAN

Generator

Random
vector z





Vector x
(vectorized image in the
range $(-1,1)/(0,1)$)

MNIST
28x28

Discriminator

Vector x




$(0,1)$

FC - fully connected layer

Sample Results - 1D-GAN to produce 2D images


Trained using MNIST digits dataset



Generated samples for different z

Not bad!

Trained using TFD (Toronto Face Dataset)




Generated samples for different z

Nearest neighbour from the training data

Sample Results - 1D-GAN to produce images

CIFAR-10 using fully-connected model for generator



Generated samples for different z

Nearest neighbour
from the training data


Not good!

DCGAN





Deep convolutional GAN

Use convolutional layers with no fully connected layers
To generate better images

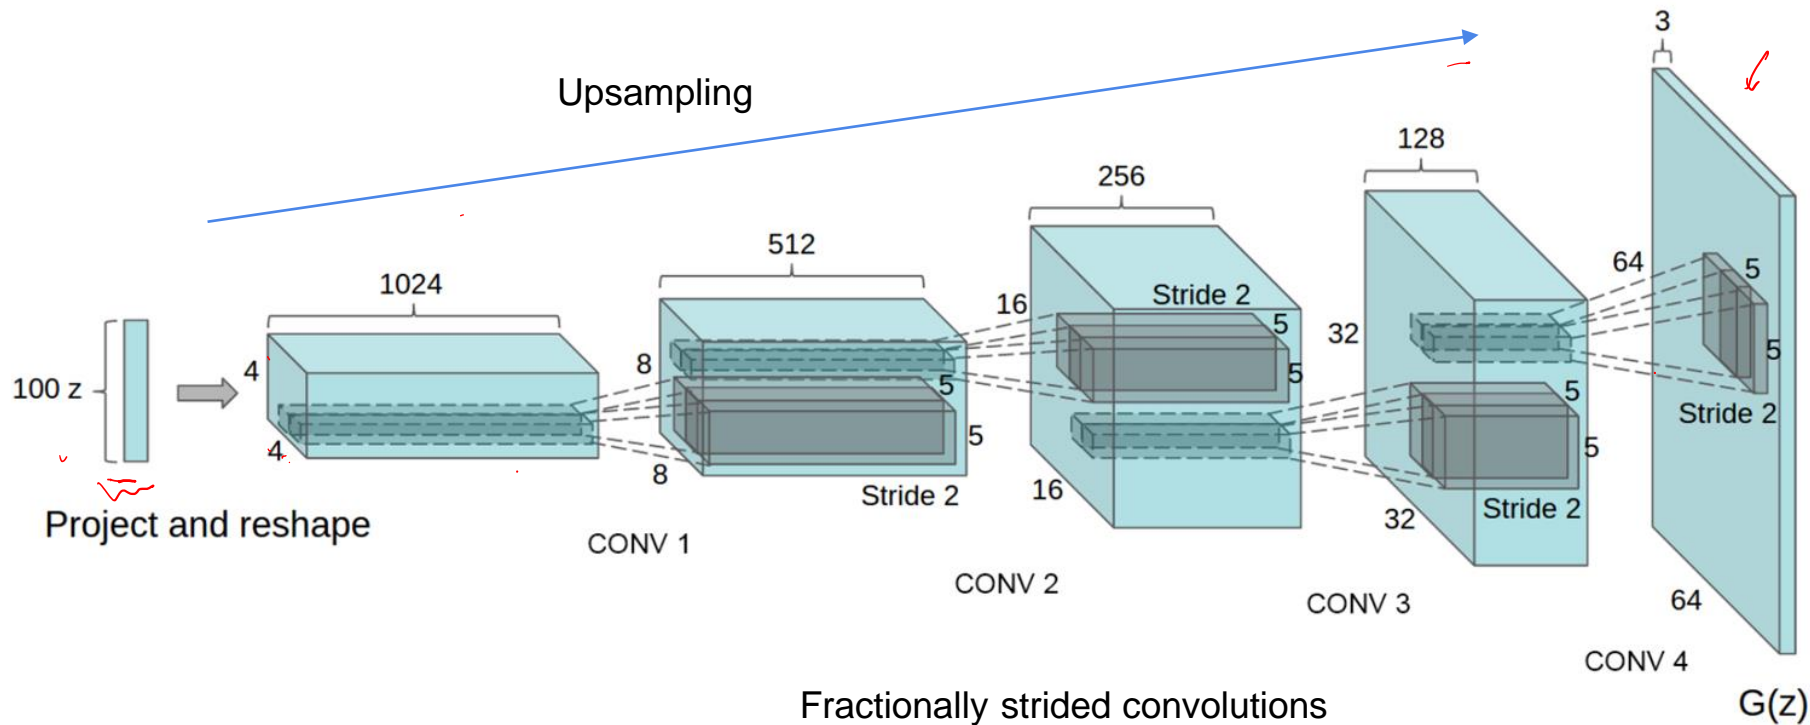
Architectural guidelines



- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove ~~fully~~ connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.



Generator in DCGAN





Generated outputs using DCGAN - a “child” writes



Groundtruth MNIST

GAN

DCGAN



These faces are
generated.
Not real!

A. Radford et al.
“Unsupervised
Representation Learning with
Deep Convolutional
Generative Adversarial
Networks” 2016

Very recent work - Generate high quality images




Timo Aila, Samuli Laine, Jaakko Lehtinen,
“Progressive Growing of GANs for Improved Quality, Stability, and Variation Tero Karras” (NVIDIA and Aalto university), submitted to ICLR 2018.
http://research.nvidia.com/publication/2017-10_Progressive-Growing-of


1024x1024 images trained using CelebA-HQ dataset

These are GAN-generated faces, and are not real.

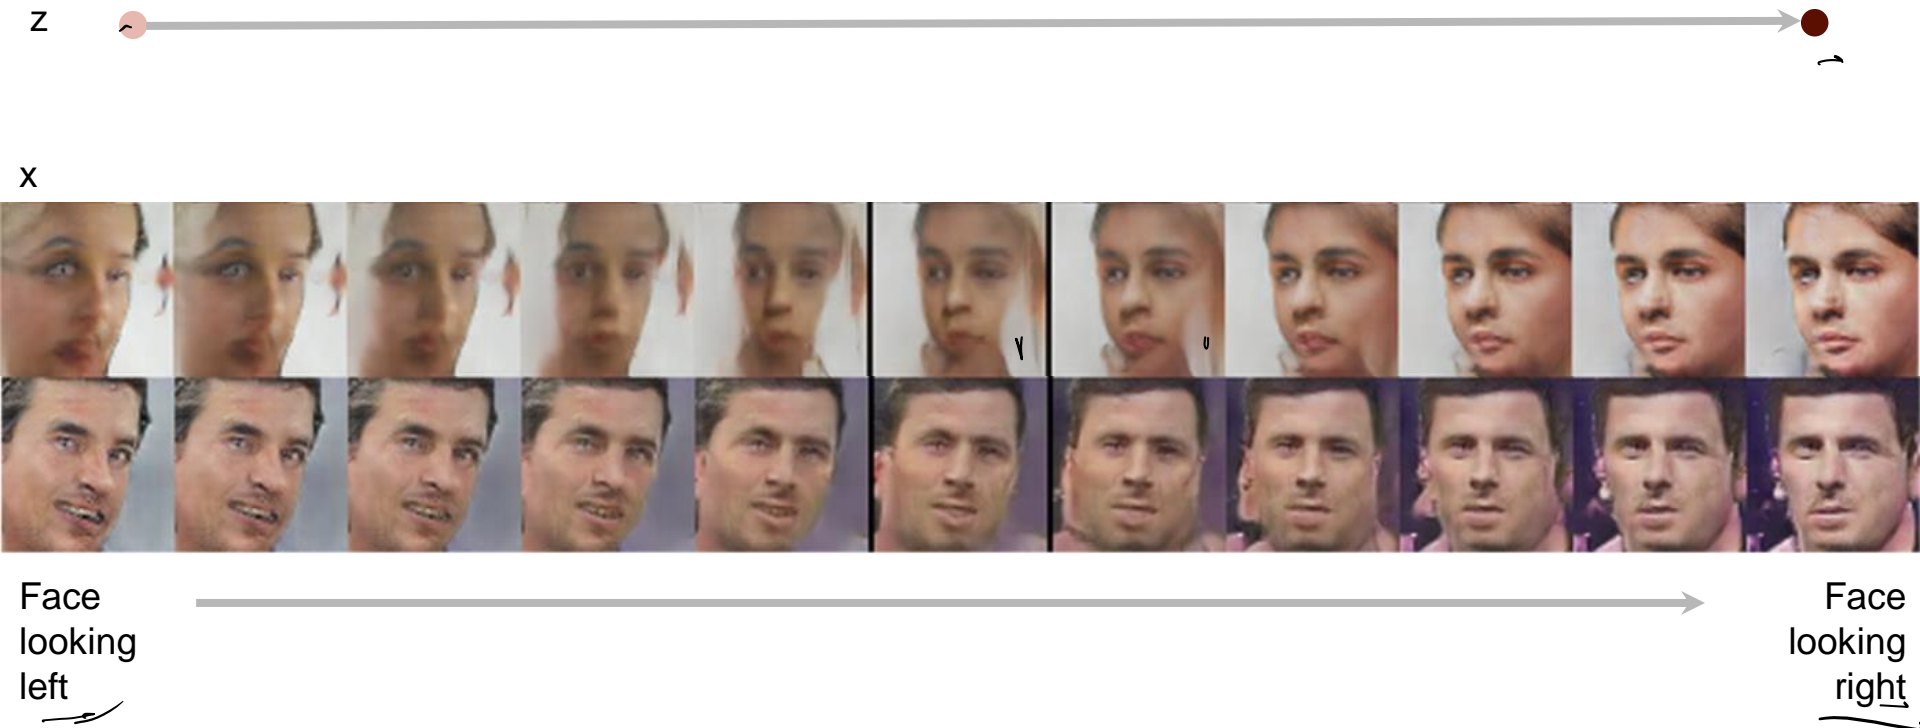
Vector interpolation for 1D GAN




Vector interpolation for DCGAN



Vector interpolation for DCGAN



Vector arithmetic in DCGAN

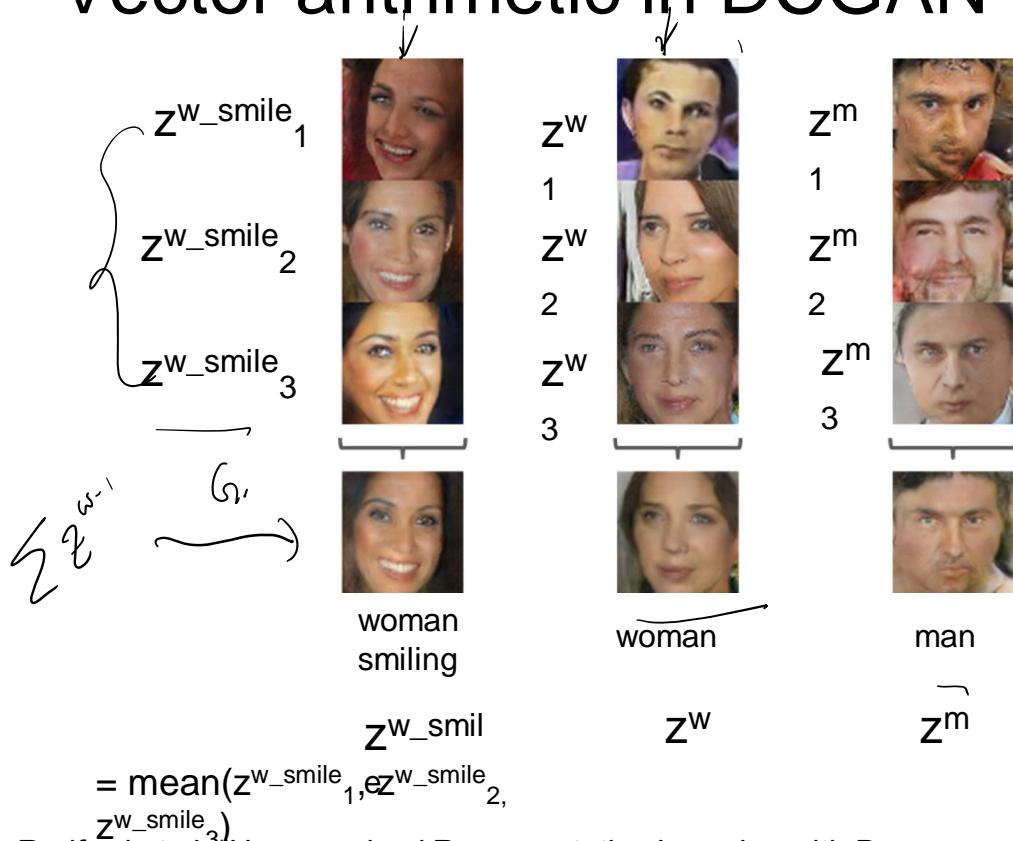


woman
smiling

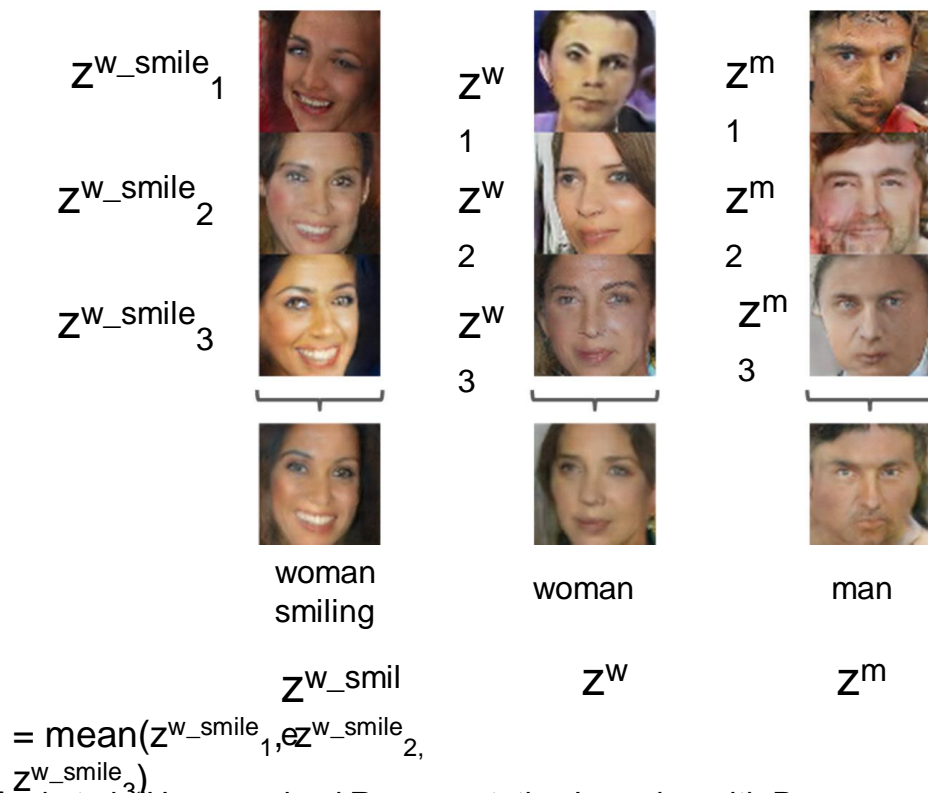
$z_{\text{w_smil}}^w$

$$= \text{mean}(z_{\text{smile}}^w, z_{\text{smile}}^w, z_{\text{smile}}^w)$$

Vector arithmetic in DCGAN

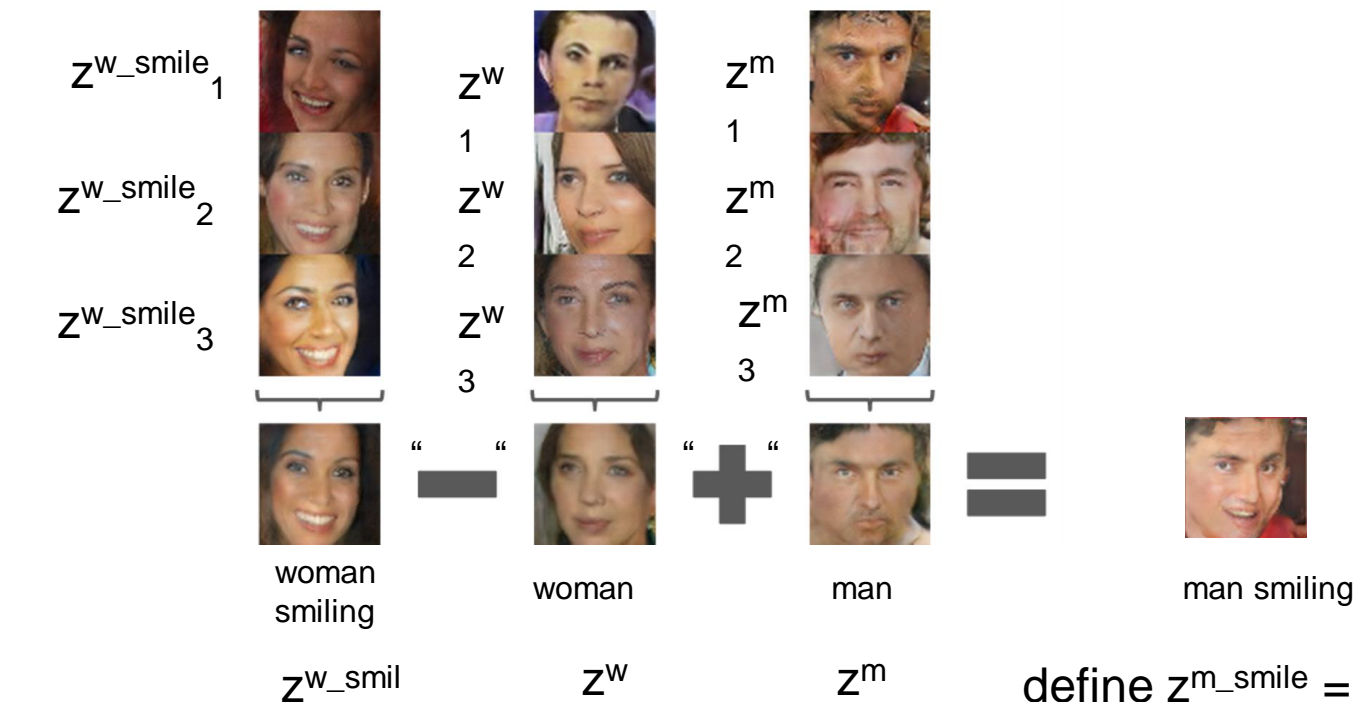


Vector arithmetic in DCGAN



$$\text{define } z^m_smile = \underline{z^w_smile - z^w} + z^m$$

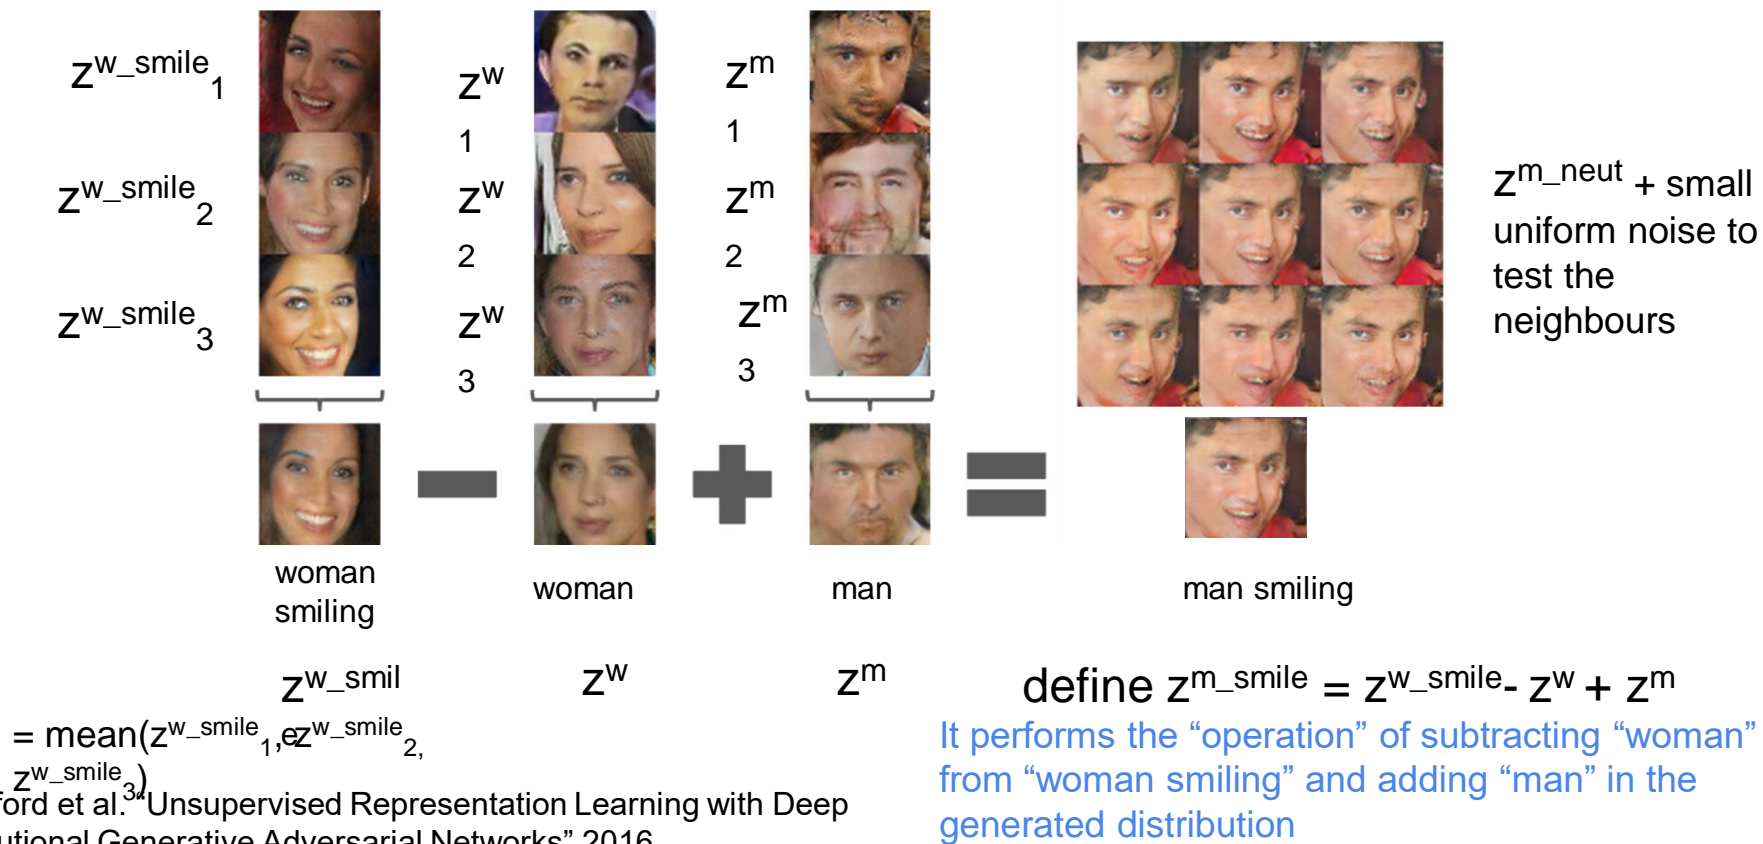
Vector arithmetic in DCGAN




define $z^m_{smile} = z^w_{smile} - z^w + z^m$

It performs the “operation” of subtracting “woman” from “woman smiling” and adding “man” in the generated distribution

Vector arithmetic in DCGAN



Vector arithmetic in DCGAN



Conditional GAN

Conditional GAN

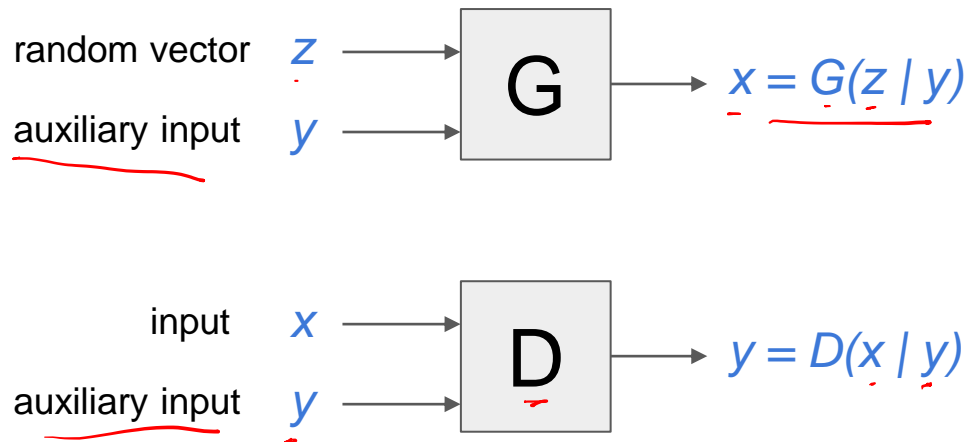
In an unconditioned generative model, there is no control on modes of the data being generated

By conditioning the model on additional information, it is possible to direct the data generation process

Generative adversarial nets can be extended to a conditional model if both the generator and discriminator are conditioned on some extra information y

y could be any kind of auxiliary information, such as class labels or data from other modalities (e.g. image)


Conditional GAN



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x} | \mathbf{y})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z} | \mathbf{y})))]$$

Conditional GAN - MNIST


Generator



The exact layers and sizes as in Mirza and Osindero “Conditional Generative Adversarial Nets” 2014 are not shown here.

Conditional GAN - MNIST


Generator



The exact layers and sizes as in Mirza and Osindero “Conditional Generative Adversarial Nets” 2014 are not shown here.


Conditional GAN - MNIST

Generator



The exact layers and sizes as in Mirza and Osindero “Conditional Generative Adversarial Nets” 2014 are not shown here.


Conditional GAN - MNIST



The exact layers and sizes as in Mirza and Osindero “Conditional Generative Adversarial Nets” 2014 are not shown here. The paper uses maxout units in discriminator.

Conditional GAN - MNIST

this row is
conditioned
on the label 3



Interesting Blogs

https://medium.com/@jonathan_hui/gan-why-it-is-so-hard-to-train-generative-adversary-networks-819a86b3750b

https://medium.com/@jonathan_hui/gan-some-cool-applications-of-gans-4c9ecca35900

<https://towardsdatascience.com/turning-fortnite-into-pubg-with-deep-learning-cycle-gan-2f9d339dcdb0>