

The background of the slide features a complex, interconnected network of white dots and lines on a dark blue background, resembling a neural network or a data mesh. The dots are of varying sizes and are connected by thin white lines, creating a dense, web-like structure that fills the entire frame.

Introduction to Neural Networks


Perceptron, XOR Problem, Multi-layer Perceptron (MLP),
Cost Functions, Activation functions, Output units

EE 5179

Instructor: Kaushik Mitra


History

- 1932 - *The Integrative Action of the Nervous System*, **Sir Charles Scott Sherrington**
 - Nervous system - interconnection of individual entities (neuron)
- 1943 - *A Logical Calculus of Ideas Immanent in Nervous Activity*, **Warren McCulloch** and **Walter Pitts**
 - McCulloch and Pitt's model
- 1949 - *The Organization of Behavior*, **Donald Hebb**
 - Hebbian learning
- 1953 - *The Perceptron*, **Rosenblatt**
- 1969 - *Limitation of Perceptrons*, **Minsky** and **Papert**
- 1980's - *Connectionism*
 - *Error Back Propagation*, **Proposed simultaneously by Many**
- 20th century Deep learning
 - CNNs - LeNet, AlexNet, VGGNet, ResNet
 - Deep LSTMs,
- The **deep saga**..... what followed is discussed in the last class




The Mark 1 Perceptron
By Rosenblatt

McCulloch - Pits model




Biological neuron




$$\sum_{i=1}^n w_i x_i > \mu$$

The input and output are binary and the weights are either excitatory (+1) or inhibitory (-1).

The Perceptron - Rosenblatt (1953)



$$\sum_{i=1}^n w_i x_i > \mu$$



*Pic courtesy, wikipedia


The Mark 1 Perceptron
By Rosenblat
for digit recognition

Weights can be any real number. He also proposed a learning algorithm for learning the weights from training data.

Perceptron - geometrical interpretation

$$\sum_{i=1}^n w_i x_i > \mu \quad , \text{What does this inequality imply in 2D case? } \text{Half plane}$$



\mathbb{X}	AND
(0, 0)	0
(0, 1)	0
(1, 0)	0
(1, 1)	1



Solve for W, μ :

$$x_1 + x_2 > 1.5$$

$$w_1 = 1, w_2 = 1 \text{ and } \mu = 1.5$$




*Pic courtesy, [cliffsnotes](#)

Any function that is linearly separable can be computed by a perceptron

Perceptron - Limitations

Goal: learn the XOR function (f^*)

\mathbf{x}	f^*
(0, 0)	0
(0, 1)	1
(1, 0)	1
(1, 1)	0



Task is adjust parameters ϑ , such that f is as close as to f^*

$$y = f(\mathbf{x}, \vartheta) \quad L = \sum_{\{\mathbf{x} \in \mathbf{X}\}} (f^*(\mathbf{x}) - f(\mathbf{x}, \vartheta))^2$$

Lets use our perceptron for f , $\vartheta = \{w, b\}$

$$f(\mathbf{x}; w, b) = w^T \mathbf{x} + b$$

Solve for $\{w, b\}$

$W = 0$, $b = 0.5$; output is 0.5 everywhere

Why this linear function can't model XOR?


How to tackle this problem?

- Can we use more than one line?
- Yes, but **how**?

Perceptron - Limitations

How to tackle this problem?

- Add a hidden layer with two units



$$y = f^{(2)}(h; U, c)$$

$$y = f^{(2)}(f^{(1)}(x))$$

$$h = f^{(1)}(x; W, b)$$


What should $f^{(1)}$ compute?

- If its linear again the composition still remains linear

$$f^{(2)}(h) = U^T h; \text{ since } h = Wx$$

$$y = U^T W x = W' x$$

Goal: learn the XOR function (f^*)




- $f^{(1)}$ should be nonlinear to extract useful features

$$h = f^{(1)}(x; W, b) = g(Wx + b)$$

- g is referred as **activation** function commonly

- We will use ReLU here


- ☐ Rectified Linear Unit (widely used)
- ☐ $g(z) = \max\{0, z\}$



Perceptron - Limitations

How to tackle this problem?

- Add a hidden layer with two units
- Use ReLU activation in 1st layer



$$y = U^T h + c; \quad y = U^T \underbrace{\max\{0, Wx+b\}}_{\text{ReLU}} + c$$


$$h = g(Wx+b)$$

Let,

$$W = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}, \quad b = \begin{bmatrix} 0 \\ -1 \end{bmatrix}, \quad U = \begin{bmatrix} 1 \\ -2 \end{bmatrix},$$

$$c = 0$$

Goal: learn the XOR function (f^*)



$$X = \begin{bmatrix} 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$

$$WX = \begin{bmatrix} 0 & 1 & 1 & 2 \\ 0 & 1 & 1 & 2 \end{bmatrix}$$

$$WX + b = \begin{bmatrix} 0 & 1 & 1 & 2 \\ -1 & 0 & 0 & 1 \end{bmatrix}$$


$$h = \begin{bmatrix} 0 & 1 & 1 & 2 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad \text{Upon ReLU}$$

After layer2

$$[0 \ 1 \ 1 \ 0]$$

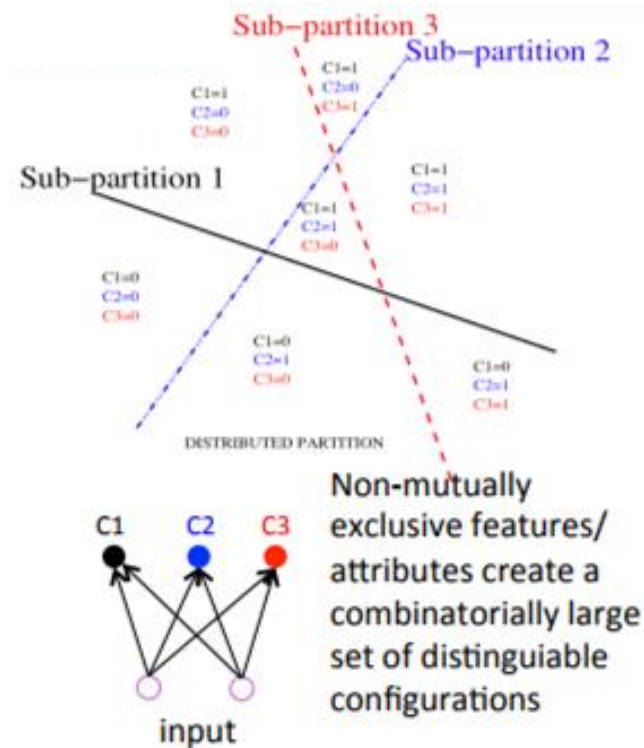
Multi-layer Perceptrons (MLP)

A typical feed forward neural network



$$\mathbf{h} = f(W\mathbf{x} + \mathbf{b}_1); \quad \mathbf{y} = g(U\mathbf{h} + \mathbf{b}_2)$$

With more hidden units network is more expressive



Feedforward Neural Networks - Cost functions

For regression,

$$J(\theta) = \frac{1}{2} \mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{\text{data}}} \|\mathbf{y} - f(\mathbf{x}; \theta)\|^2$$

$$\frac{1}{2} \sum_{\{x_i, y_i\}} \|y_i - f(x_i, \theta)\|^2$$

For classification,

- Typically outputs a probability vector $q(c = k | x) \forall k$
- How do you compare two distributions?
 - KL divergence, $KL(p||q)$


$$D_{KL}(p(x)||q(x)) = \sum_{x \in X} p(x) \ln \frac{p(x)}{q(x)}$$

$$= \sum p(x) \ln p(x) - p(x) \ln q(x)$$

$$= \underbrace{-H(p)}_{\text{entropy}} + \underbrace{H(p, q)}_{\text{cross-entropy}}$$

$$J(\theta) = \sum_{x_i, y_i} H(p(x_i), q(x_i))$$


Activation functions




$h = g(Wx+b)$; Affine transformation followed by activation function, g

- Very important factor in learning features


$$g(z) = \max\{0, z\}, \text{ ReLU}$$




$$g(z) = \sigma(z), \text{ sigmoid}$$



$$\tanh(z) = 2\sigma(2z) - 1$$



Output units




- Linear units for real valued outputs
 - ❑ Activation function is left to be linear
 - ❑ Given features h ,
$$\mathbf{y}' = \mathbf{W}\mathbf{h} + \mathbf{b}$$
 - ❑ Most commonly used with regression tasks

- Say you want to do binary classification
 - ❑ What kind of distribution describes output?
Bernouli
 - ❑ How to constrain the output - valid probability?
Can you use linear activation?
 - $$P(y = 1 \mid \mathbf{x}) = \max \left\{ 0, \min \left\{ 1, \mathbf{w}^\top \mathbf{h} + b \right\} \right\}.$$
 - ❑ What is the problem? *Not amenable for gradient based learning*
 - ❑ Instead, use sigmoid unit - output $\in [0,1]$

$$\hat{y} = \sigma(\mathbf{w}^\top \mathbf{h} + b)$$

Output units



- Now, say we want to do multi-class classification (K classes)

- ❑ Output should be K probabilities,
 $p_k = p(\text{class} = k | \mathbf{x}) \quad \forall k = 1 \text{ to } K$

- ❑ Can we use K sigmoid units?

Won't be sufficient, since probabilities are not constrained to sum to 1

$$\sum_k p_k = 1$$

- ❑ We will look at **softmax** unit for this

Idea is to convert a vector of real values to valid probabilities,

Hqw?

- ❑ Make all the elements positive

- ❑ Normalize the values

- Let, $\mathbf{z} = [z_1, \dots, z_K]^T$; $\mathbf{z} = \mathbf{W}\mathbf{h} + \mathbf{b}$

$$\text{softmax}(\mathbf{z})_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}.$$